

Models to run (ordered):

Note: we will run all models in their base version, with 60 frames, fps = 4, for 15 epochs

Model	Encoder command	Status	Submitted by	Best Test MSE	Notes
DINOv2 small	--dino_v2_small	Finished	Aicha	Lr = 0.001 Wd = 0.8 MSE = 0.0055	
DINOv2 base	--dino_v2_base	Finished	Lynn	Lr = 0.001 Wd = 0.8 MSE = 0.0052	
DINOv2 large	--dino_v2	Finished	Aicha	Lr = 0.001 Wd = 0.8 MSE = 0.005	Waiting time has been of 2 days. It only started running around 5 a.m. on Wednesday morning
DINOv2 giant	--dino_v2_giant	Finished	Lynn	Lr = 0.001 Wd = 0.01 MSE = 0.0051	Had to reduce the batch size from 10 to 6
vJEPa 2	--vjepa2	Finished	Lynn	Lr = 0.001 Wd = 0.01 MSE = 0.0055	
vJEPa2 on shuffled frames	--vjepa2		Aicha		
VideoMAE	--videomae_b	Run again	Aicha		Concatenate frames s.t. the model sees all the frames
VideoMAE V2	--videomae_v2_b	-	Aicha		Will not run this model
DINOv3	--dino_v3	Finished	Lynn	Lr = 0.001 Wd = 0.01 MSE = 0.0048	

DINOv2 robust small	--dino_v2_robust_small	Running	Aicha	Lr = 0.001 Wd = 0.01 MSE = 0.0058	Tested! All good
DINOv2 robust base	--dino_v2_robust_base	Running	Aicha	Lr = 0.001 Wd = 0.8 MSE = 0.0053	
DINOv2 robust large	--dino_v2_robust_large	Finished	Aicha	Lr = 0.001 Wd = 0.8 MSE = 0.0051	
VideoMamba (implement)	--				If I have time, copy the env (transformers_v2)
ConvNeXt	--convnext	Finished	Lynn	Lr = 0.001 Wd = 0.01 MSE = 0.0074	
ConvNeXt V2	--convnext_v2	Paused	Aicha		
SigLIP2 (test)	--siglip2		Aicha		
ViT	--ViT	Finished	Lynn	Lr = 0.001 Wd = 0.8 MSE = 0.0066	
ResNet-50	--resnet-50		Lynn		
ResNet-50 Linf (implement + test)	--		Aicha		
AlexNet	--alexnet	Paused	Lynn		
iJEPA (not sure we will run this)	--				
X3D	--x3d_m				Concatenation as videomae
SlowFast (not implemented yet)	--				